

Gør tanke til handling
VIA University College



Implementing an object model on a relational database

Second Semester Project (SEP2), Fall 2018
Bo Brunsgaard, BOBR@VIADK

VIA University College

2 november 2018

1

PLEASE NOTE:

This presentation was intended to be used as background materiale for
a step-by-step case, done on the black board

If life's circumstances conspired to keep you absent from the SEP2 session on
Nov. 1st, the contents here may be a bit hit-and-miss.

Life sucks that way.

VIA University College

2 november 2018

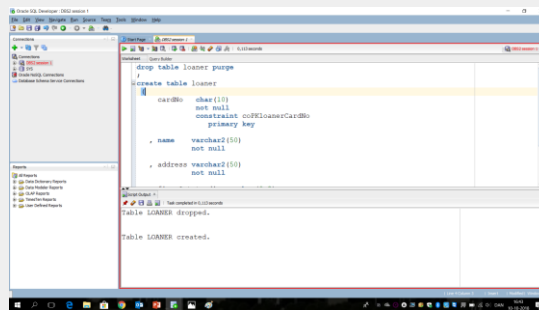
2

Agenda

- OO, ER and relational: same same, but different
- Case walk-through

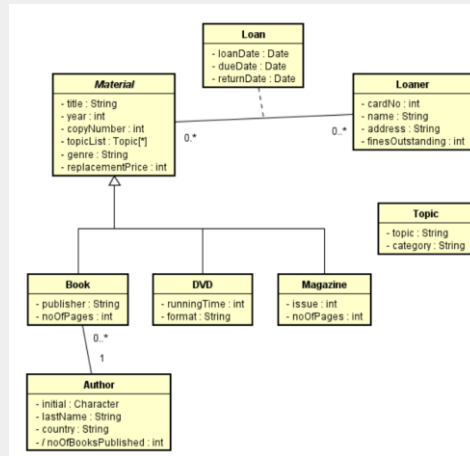
A note on tools

- You are using postgresSQL
- My examples are built using the Oracle RDBMS
- Tools *look* different
 - Mapping principles are exactly the same
 - SQL is (almost) the same
- I have certain habits – e.g. always naming my constraints
 - You may do things differently in DBS1; do as your DBS1 teacher says



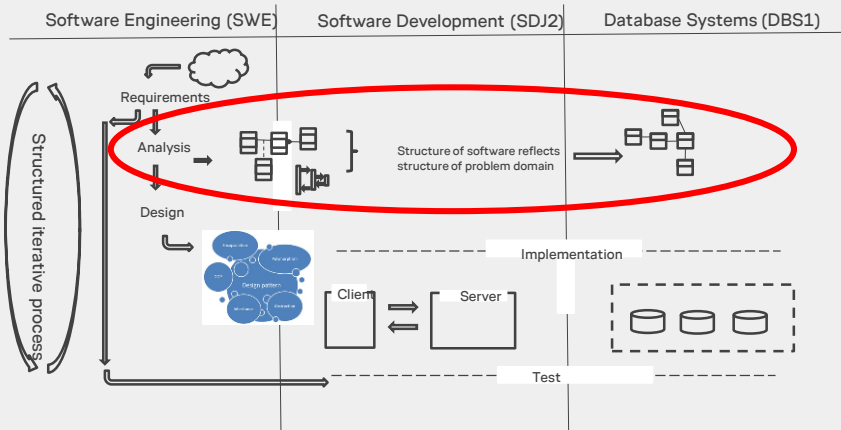
A note on examples

- We will be using a small (analysis) class model for a minimal library system
- In some of the examples, I may make small deviations from the model, either to demonstrate a point or to avoid clutter.



First principles

Remember this?



VIA University College

02-11-2018

7

First principles

1. The structure of both the software and the database reflect the structure of the problem domain
2. The structure of the software and the database should, ideally, be the same
 - For technical reasons they cannot be exactly the same (impedance mismatch)
3. We capture the problem domain structure in a domain (analysis) model

Collary:

1. The database model & structure is DERIVED from the (object-oriented) domain model
2. Analysis/ domain modelling and database design are NOT disjoined activities!

VIA University College

2 november 2018

8

Concepts

Book
- title : String
- year : int
- author : String
- noOfPages : int
- onLoan : boolean
+ Book() : Book
+ getTitle() : String
+ getYear() : int
+ getAuthor() : String
+ getNoOfPages() : int
+ loan() : void
+ returnLoan() : void



Book
title
year
author
noOfPages
onLoan



```
create table book
(
  title  varchar2(50)
        not null

, year   number(4,0)
        not null
        constraint coCHBookYear
        check (not year < 1438 )

, author varchar2(50)
        not null

, noOfPages number(4,0)
        not null
        constraint coCHBookNoOfPages
        check (noOfPages > 0 )

, onLoan char(01)
        not null
        constraint coCHBookonLoan
        check (onLoan in ('Y','N') )

, constraint coPKBook
  primary key (title, year, author)
)
;
```

- Conceptually, class, entity and table are closely related.
 - An entity is the class without operations
 - The table is an implementation of an entity

VIA University College

2 november 2018

9

First-cut mapping

- A class becomes an entity / table
 - Collary: an object becomes a row of said table
- An attribute of a class becomes a column of a table
- An association between two classes becomes a primary/foreign key relationship
 - Collary: we need to identify which attributes of a class make up the foreign key!
- An association class becomes a table, holding the PKs of the tables at either end
- Specialized classes (can.) become separate entities/tables
 - Each table must then have all columns inherited from the generalized class
- Abstract classes?
 - Abstract classes? We don't have no stinking abstract classes!

VIA University College

2 november 2018

10

Object-relational mapping



- You will need to code to:

- translate from four tables into one object on reading a book from the database
- Translate from one object into on four tables on inserting/updating/deleting

Object-Oriented models, Entity-Relation models, and relational implementation

At first cut, transforming a class model to an entity model, appears straight-forward



The devil is in the detail:

- Identity vs. Data values
- Embedding vs. key value relationships
- Complex objects vs. simple tables
- Derived attributes
- Object-at-a-time data redundancy vs. normalized relations
- Specialization and inheritance

Gør tanke til handling
VIA University College

Identity



Identity: theoretical concerns

Identity vs. data values

- For an object, state (field values) and identity are two different things

```
private Book book1 = new Book("Das Kapital", 1867, "K. Marx");  
private Book book2 = new Book("Das Kapital", 1867, "K. Marx");
```

- We reference them by the field that holds the pointer to the object

Identity vs. data values

```
create table book
(
  title      varchar2(50)
            not null
, year      number(4,0)
            not null

, authorInit char(01)
            not null

, authorLastName varchar2(50)
            not null

, noOfPages  number(4,0)
            not null

, constraint coPKBook
  primary key (title, year, authorInit, authorLastName)
```

- For a row in a table, identity is given by state
- We must identify the set of fields from the class definition that will uniquely identify an object/row -> primary key

Identity vs. data values

- An object has identity and state
- Two objects with exactly the same data will be two separate objects

```
private Book book1 = new Book("Das Kapital", 1867, "K. Marx");
private Book book2 = new Book("Das Kapital", 1867, "K. Marx");
```

Identity vs. data values

- An object has identity and state,
- A row in a table has only state: identity is given by state
- Two rows in a table with the same data will be duplicates:

```
insert into book (title, year, author, noOfPages, onLoan)
values ('Das Kapital', 1867, 'K. Marx', 1168, 'N')
;
insert into book (title, year, author, noOfPages, onLoan)
values ('Das Kapital', 1867, 'K. Marx', 1168, 'N')
;
```

Script Output x | Task completed in 2,142 seconds

Error report -
 SQL Error: ORA-00001: entydig begrænsning (DBS2.COPKBOOK) er overtrådt
 00001, 00000 - "unique constraint (%s.%s) violated"
 *Cause: An UPDATE or INSERT statement attempted to insert a duplicate key.
 For Trusted Oracle configured in DBMS_MSC mode, you may see

Identities



- Theoretically, this is a big deal. In practice, not so much
- Usually it indicates sloppy analysis and a bad domain modeling!
 - Often the result of mixing analysis and design classes
 - Real-world (analysis) objects can normally be told apart by their values
 - Design objects (e.g. connections, UI elements etc.) often cannot
- Think about it: in which system would we have two different students with the same studentID, social security number, name, gender, address, etc?

Identity: practical concerns

ID column

- IF the difference between state and identity is important, add an artificial key, and designate that as `primary key`



```
create table book
(
  id          number(10,0)
             generated always as identity
             start with 1 increment by 1 noMaxValue


             constraint coPKBook primary key

, title      varchar2(50)
             not null

, year       number(4,0)
             not null
             constraint coCHBookYear
```

ID and natural key

(one of my hobby horses Rant warning!)



- If you add an artificial key, the business key uniqueness rule still exists and must be enforced!

```
create table book
(
  id          number(10,0)
             generated always as identity
             start with 1 increment by 1 noMaxValue
             constraint coPKBook primary key
  , title     varchar2(50)
  , year      number(4,0)
  , authorInit char(01)
  , authorLastName varchar2(50)

  , constraint coUNBookBusinessKey unique
    (title, year, authorInit, authorLastname)
```

VIA University College

2 november 2018
23

Lame uses of ID


(one of my hobby horses Rant warning!)



- It has become customary always to add an artificial key to tables
 - (I am fighting a losing battle against this)

But this does SO not
fix any normalization
issues

Even
if:



VIA University College


2 november 2018
24

Lame uses of ID: normalization

(one of my hobby horses Rant warning!)



- Normalization identifies functional dependencies on the primary key
- Some believe that an artificial ID makes the primary key go away, and tables automatically on third normal form



VIA University College

2 november 2018


26

Lame uses of ID: normalization

(one of my hobby horses Rant warning!)



- Adding an artificial key just moves normalization problems one step



- Table has a partial functional dependency, and is thus not on 2NF

VIA University College

2 november 2018

26


Lame uses of ID: normalization

(one of my hobby horses. Rant warning!)



- Adding an artificial key just moves normalization problems one step

Artificial primary key



- Table now has a single-column PK, and is thus trivially on 2NF
 - A value cannot be determined by a sub-part of the PK if the PK has only one part
- But it now has a transitive dependency, ID -> movieTitle -> runningTime
- So we just exchanged a **1NF problem** for a **2NF problem**

VIA University College

2. november 2018

27

Gør tanke til handling
VIA University College

Relationships

VIA University College

2. november 2018

28

Pointers vs. Shared data values


- Objects reference other objects by pointers

```
public class Loan {
    private book    bookLoaned
    private Loaner  loaner
    private Date    dateloaned
}
```

- Table rows reference other rows by keeping a copy of data values of key columns (primary/foreign key relationship)

```
create table loan
(
    title    varchar2(50)
  , year    decimal(4,0)
  , .....
  , constraint coFKLoanMaterial
    foreign key (title, year) references book
)
```

From reference to shared values (1 of 2)



```
public class Author
{
    private String initial;
    private String lastName;
    private String country;
}
```

```
public class Book
{
    private String title;
    private int year;
    private Author author;
}
```

- Relationships between objects are done by embedding a reference to an instance of another class
- In a relational database, relationships are implemented by common data values

From reference to shared values (2 of 2)

```
create table author
(
  init      char(02)
           not null
, lastName  varchar2(50)
           not null
, constraint coPKAuthor
  primary key (init, lastName)
);
```


```
create table book
(
  title      varchar2(50)
           not null
, year       number(4,0)
           not null
, constraint coCHBookYear
  check (not year < 1438 )
, authorInit char(01)
           null
, authorLastName varchar2(50)
           null
, constraint coFKBookAuthor
  foreign key (authorInit, authorLastName) references author
);
```

- Relationships between objects are done by embedding an instance of another class
- In a relational database, relationships are implemented by common data values

Multiplicities: 1-N

- Once we have identified primary keys, mapping associations to PK/FK relationships is usually straight-forward

One-to-many:



```
create table book (
  title      varchar2(50)
, year       decimal(4,0)
, authorInit char(01)
           not null
, authorLastName varchar2(50)
           not null
...
, constraint coFKBookAuthor
  foreign key (authorInit, authorLastName)
  references author
);
```


Multiplicities: N-M

- Once we have identified primary keys, mapping associations to PK/FK relationships is usually straight-forward

many-to-many:



Create a new table for the relationship, holding the PKs of the classes on either side

```
create table bookTopic
(
    title    varchar2(50)
  , year    decimal(4,0)
  , topic   varchar2(50)

  , constraint coPKBookTopic
    primary key (title,year,topic)
  , constraint coPKBookTopicBook
    foreign key (title,year) references book
  , constraint coPKBookTopicTopic
    foreign key (topic) references topic
)
```

- Note the composite PK of the relationship table
- Note the two FK constraints, pointing to each their side

VIA University College

2 november 2018
33

Multiplicities: 1-1

- Once we have identified primary keys, mapping associations to PK/FK relationships is usually straight-forward

one-to-one:



Copy the columns of the PK from one side as a FK on the other side. Make the FK unique.

```
create table emperor
(
    name        varchar2(50)
    constraint coPKEmperor
    primary key

  , empire      varchar2(50)
    constraint coFKEmperorEmpire
    references empire
    constraint coUNEEmperorEmpire
    unique
)
```

- The FK can be placed on either side of the relationship, unlike in a one-to-many implementation
- The foreign key is declared UNIQUE to ensure the 1-1 property

VIA University College

2 november 2018
34


Relationships

Subtleties

Direction of relationships (navigatibility)

- Conceptually, associations are UNI-directional
 - PK/FK relationships are always BI-directional
 - Primarily a theoretical issue more than a practical one
 - It can have an effect on:
- Applies regardless of multiplicity of relationships

Directed one-to-one:



Copy the columns of the PK from the TO-side as a FK on the FROM side

```

create table emperor
(
    name          varchar2(50)
    constraint coFKEmperor
    primary key

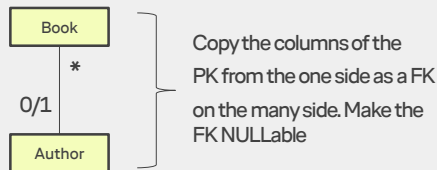
, empire         varchar2(50)
    constraint coFKEmperorEmpire
    references emperor
    constraint coUNEmperorEmpire
    unique
)
  
```

- (it actually won't make any difference to the database (still bi-directional), but matches the code structure better)

Optional relationships

- An optional relationship is still implemented as a foreign key
- To make it optional, it is declared as `Nullable`
 - A foreign key constraint isn't checked if the actual value of the columns is/are `NULL`

Optional one-to-many



```
create table book (
    title    varchar2(50)
  , year    decimal(4,0)
  , authorInit char(01)
    null
  , authorLastName varchar2(50)
    null
  ...
  , constraint coFKBookAuthor
    foreign key (authorInit, authorLastName)
    references author
```


VIA University College

2 november 2018

37

Aggregation

- In a composition, object the subordinate class **MUST** be related to an object of the superior class, **AND** the superior class is responsible for creating and deleting subordinate objects



```
create table book (
    title    varchar2(50)
  , year    decimal(4,0)
  , seriesName varchar2(50)
    not null
  , constraint coFKBookSeries
    references series
    on delete cascade
  ...
```

- Note that the foreign key is **NOT NULL**, since a book **MUST** be part of a series

VIA University College

2 november 2018

38

Complex objects and simple tables

Complex objects – simple tables


Class/object

- Complex data types -> a field can hold a reference to a complex object
- Multiple values as array, arrayList, hashMaps..
- Arrays (etc) can again hold references to complex objects

Entity/table

- Columns of primitive type (date is kind-of a bit of an exception)
- Have single-valued columns (i.e. no repeating groups)

Complex objects




author is an instance of the Author class
which has many fields

topics is an array. A single book can have
zero, one or many topics

- Objects can have a complex structure, e.g.:
 - Embedding objects (with multiple fields) of other classes
 - A field can contain multiple values (as arrays, ArrayLists etc.)

Complex objects: objects references




```
create table book
(
  title      varchar2(50)
            not null
, year      number(4,0)
            not null
, authorInit char(01)
            not null
, authorLastName varchar2(50)
            not null
, constraint coFKBookAuthor
  foreign key (authorInit, authorLastName)
  references author
```


- A field of type Class-something must be:
 - Created as a new table, representing Class-something
 - replaced by a foreign key to the table representing it

Object references (minor rant)

- Anyway, to me this is sloppy documentation practice



- It means exactly the same as this
- Much easier to understand at a glance




VIA University College

2 november 2018

43

Complex objects: arrays



```

create table bookTopic
(
    bookTitle      varchar2(50) not null
    , bookYear      number(4,0)  not null
    , authorInit    char(01)     not null
    , authorLastName varchar2(50) not null
    , topic         varchar2(50) not null

    , constraint coPKbookTopic
      primary key (bookTitle, bookYear, authorInit, authorLastName, topic)

    , constraint coFKbookTopic
      foreign key (bookTitle, bookYear, authorInit, authorLastName)
        references book
)
  
```

- A multi-valued field must be split into a separate table.
- This will represent a 1 . . * relationship (which is what the array actually is - one book contains (has) many topics)

VIA University College

2 november 2018

44

Complex objects: arrays

Book
- title : String
- year : int
- author : Author
- genre : String
- replacementPrice : int
- topics : String[]

```
create table bookTopic
(
  bookTitle      varchar2(50) not null
, bookYear       number(4,0)  not null
, authorInit     char(01)     not null
, authorLastName varchar2(50)  not null
, topic          varchar2(50) not null

, constraint coPKbookTopic
  primary key (bookTitle, bookYear, authorInit, authorLastName, topic)

, constraint coFKbookTopic
  foreign key (bookTitle, bookYear, authorInit, authorLastName)
    references book
)
```

- The new table must have a foreign key to the original table (the new table is the many-side of the relationship)
- If the array was of a primitive type, PK = original table's PK + the value from the array

VIA University College

2 november 2018

45

Complex objects: arrays (a probably better solution in most cases)

- Often better to create a table for the values that can appear in the array
 - All books now use same set of values
 - It now becomes a ** relationship.

```
create table topic
(
  topic varchar2(50) not null
  constraint coPKtopic
    primary key
)
;
```

```
create table bookTopic
(
  bookTitle      varchar2(50) not null
, bookYear       number(4,0)  not null
, authorInit     char(01)     not null
, authorLastName varchar2(50)  not null
, topic          varchar2(50) not null
  constraint coFKbookTopic
    references topic

, constraint coPKbookTopic
  primary key (bookTitle, bookYear, authorInit, authorLastName, topic)

, constraint coFKbookTopicbook
  foreign key (bookTitle, bookYear, authorInit, authorLastName)
    references book
)
```

VIA University College

2 november 2018

46

NO. Just f* NO



```
insert
into book (title, year, authorInit, authorLastName, topics)|
values ( 'It can''t happen here'
      , 1935
      , 'S'
      , 'Lewis'
      , 'politics, populism, totalitarianism, dystopian fiction, democracy'
    )
```

- The solution is NOT to store a comma-separated list in a string
 - Unsearchable, unindexable, unconstrainable, un-everything

VIA University College

2 november 2018

47

Handling complex objects - follow-on effects on code



- You will need to code to:
 - translate from four tables into one object on reading a book from the database
 - Translate from one object into on four tables on inserting/updating/deleting

VIA University College

2 november 2018

48

O/R mapping, part 1

```

preparedStatement stmtBook = conn.prepareStatement("select b.title, b.year, a.initial, a.lastName, a.country "
+ " , b.genre, b.replacementPrice "
+ "from book b "
+ " join author a "
+ " on (      a.init = b.authorInit "
+ "        and a.lastName = b.authorLastName "
+ "        )"
+ "where b.title = ? "
+ " and b.year = ? "
+ " and b.authorInit = ? "
+ " and b.authorLastName = ? ")

...

stmtBook.setString("title",searchParam.getTitle())
... ..
resultSet rsBook = stmtBook.executeQuery()

if (rsBook.next()) {
    this.title = rsBook.getString(1)
    ...
    this.replacementPrice = rsBook.getInt(7)
}

```

Book
- title : String
- year : int
- author : Author
- genre : String
- replacementPrice : int
- topics : String[*]

- Reading a book object: a JOIN between book and author

VIA University College

2 november 2018

49

O/R mapping, part deux

```

preparedStatement stmtTopics = conn.prepareStatement("select t.topic
+ "from bookTopic bt "
+ " join topic t "
+ " on (      t.topic = bt.topic )"
+ "where bt.title = ? "
+ " and bt.year = ? "
+ " and bt.authorInit = ? "
+ " and bt.authorLastName = ? ")

stmtBook.setString("title",searchParam.getTitle())
... ..
resultSet rsTopics = stmtTopics.executeQuery()

while (rsTopics.next()) {
    this.topics.add(rsTopics.getString(1))
}

```

Book
- title : String
- year : int
- author : Author
- genre : String
- replacementPrice : int
- topics : String[*]

- Entries in the topic array must be read and the array populated with all returned rows


VIA University College

2 november 2018

50

Derived attributes

When an attribute really isn't there?



- `noOfBooksPublished` is not a field *per se*:
- It is derived from other information
 - In this case from the cardinality of the association with the `Book` class
- (In the design model it might very well be transformed into an operation (method) instead)

Derived attributes

- Storing the derived value negates the idea of a derived field
- Leading to problems keeping it up-to-date
 - When a new book is added, we must remember to update the author
 - When a book is deleted, we must remember to update the author
 - When a book is updated, worst case, we must remember to update TWO authors
- So: derived attributes should NOT be mapped to columns of a table

Derived attributes: views:

```
create view author as
select a.init
      , a.lastName
      , a.country
      , count(*) as noOfBooksPublished
from   authorBaseTable a
join   book b
on     ( b.authorInit    = a.init
        and b.authorLastname = a.lastName)
group by a.init
        , a.lastName
        , a.country
;
```

- We can create a view that calculates (i.e. derives) the value on-the-fly

Derived attributes: trouble with (this) view:

```
create view author as
select a.init
      , a.lastName
      , a.country
      , count(*) as noOfBooksPublished
from   authorBaseTable a
join   book b
on     ( b.authorInit = a.init
        and b.authorLastname = a.lastName)
group by a.init
        , a.lastName
        , a.country
;
```

- The view requires us to join author with book, to get number of books
- Now it cannot be updated!
 - Not that we would need update noOfBooksPublished anyway, but the join gets us

Derived attributes: a possible solution

```
create view author as
select a.init
      , a.lastName
      , a.country
      , count(*) as noOfBooksPublished
from   authorBaseTable a
join   book b
on     ( b.authorInit = a.init
        and b.authorLastname = a.lastName)
group by a.init
        , a.lastName
        , a.country
;
```

```
create or replace trigger trISUPDAuthor
instead of update on author
for each row
begin
update authorBase
  set init      = new.init
    , lastName = new.lastName
    , country  = new.country
  where init    = old.init
    and lastName = old.lastName
;
end
;
```

- We can use an `instead of trigger` to capture the update, and transform it into just an update on the `author` table (disregarding the association and join to `book`)
- Of course, we need to do something similar in case of `insert` and `delete`

Derived attributes (minor rant)



- Dealing with views and instead of triggers can become very convoluted, with little benefit
- Usually, I wouldn't bother handling derived attributes in the database
 - Let the application code handle it – calculate it in the `getNoOfBooksPublished()` method on the `book` object
- Exceptions may apply with the database is used by other applications and access through other interfaces

VIA University College

2 november 2018
57

Gør tanke til handling
VIA University College

Normalization

**THIS "NORMAL"
YOU SPEAK OF
DOESN'T SOUND
FUN AT ALL**

VIA University College

2 november 2018
58

Object-at-a-time vs. Normalization

(my take, others may disagree!)


Book
- title : String
- year : int
- author : Author
- copyNumber : int
- genre : String
- replacementPrice : int
- noOfPages : int
- onLoan : boolean

- Classes are often designed with an eye towards the specific application's needs
- Applications deal with individual objects, and don't consider objects that they are not dealing with
- So, it makes sense that if I look at a particular copy of a book, I get all the information about it (as above)

VIA University College

2 november 2018
59

Table design from class:



- Table is only on 1NF (it has a partial functional dependency)
- `onLoan` depends on the entire primary key. You borrow a specific copy of a book
- `noOfPages`, `genre` and `replacementPrice` however, depend only on part of the primary key. They are the same for all copies of a book

VIA University College

2 november 2018
60

Table design from class:

- We have both INSERT, UPDATE, and DELETION anomalies
 - Your friendly neighborhood DBS1 teacher will supply all the details on this
- For now, suffices to say



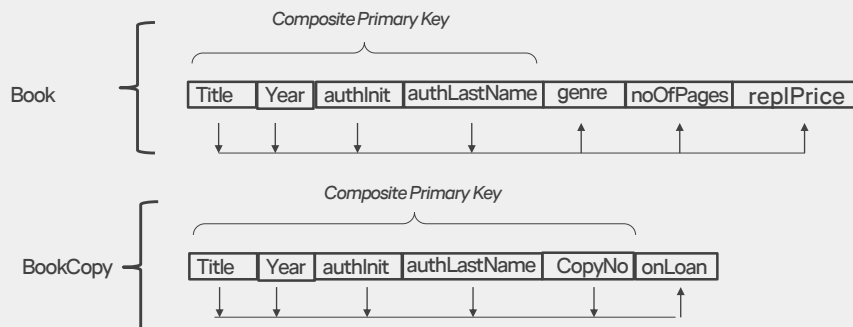
VIA University College

2 november 2018
61

Object-at-a-time vs. Normalization

(my take, others may disagree!)

- We need to split the table into two




VIA University College

2 november 2018
62

Specialization and inheritance

Specialization and inheritance

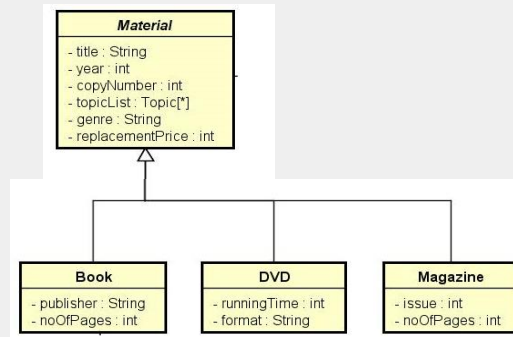
- Specialization and inheritance are elegant ways of handling common/separate properties



- In the library, we have books, DVDs and magazines
- They are all different, with each their attributes

Specialization and inheritance

- But books, DVDs and Magazines have attributes in common
- We model these as a common ancestor, and books, DVDs and magazines inherit attributes from this



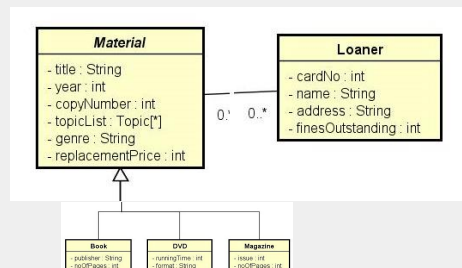
VIA University College

2 november 2018

65

Specialization and inheritance

- Even better: shared relationships can be made to the super-class
- I simply loan a material which is then a book, a DVD or a magazine
- When I need it to be a material, it will be a material
- When I need it to be a book, it will be a book



VIA University College

2 november 2018

66

Specialization and the relational model

- The relational model does not support concepts of generalization, specialization and inheritance
 - OK, there's something called object-relational databases. Avoid like the plague! Just my opinion.
- We need to map to independent tables
 - Generalization and specialization in individual tables?
 - Only specializations?
 - Merge generalization and specializations into a single table?
- There is no single "right" solution!
- Let's do the rest on the blackboard